# Siva Krishna Reddy Kunchala

Bloomfield Hills, MI (Open to Remote or Relocation) | +1 (660) 270-6448 | Sivakrishnareddykunchala18@gmail.com | LinkedIn | GitHub | Portfolio

## PROFESSIONAL SUMMARY

AI Engineer with 3.5+ years across production software, cloud operations, and LLM-based automation. I ship agentic workflows that call REST APIs to validate data and take actions in enterprise and supply-chain systems, using structured outputs, deterministic validation, exception routing, and human-in-the-loop review. Built evaluation pipelines that measure every retrieval change. I decide when a step needs an LLM, rules-based code, or a human, and explain those tradeoffs to operations and business stakeholders. Built and run a production serverless AI agent on AWS (Lambda, Bedrock, DynamoDB), backed by hands-on AWS operations experience.

## TECHNICAL SKILLS

**Languages:** Python, SQL, Bash
**LLM & Agentic AI:** Azure OpenAI, Amazon Bedrock, LangGraph, LangChain, Tool/Function Calling, Structured Outputs (Pydantic), Multi-Agent Orchestration, Prompt Engineering, Human-in-the-Loop Workflows
**Evaluation & Reliability:** Retrieval Evaluation (Recall@5, MRR), Checkpointing, Retries and Failure Recovery, Exception Classification
**Integration:** REST APIs, HubSpot CRM, JSON, Data Reconciliation, ETL Pipelines
**Retrieval (RAG):** FAISS, Sentence Transformers, Hybrid Search, Reranking
**Cloud & DevOps:** AWS (Lambda, API Gateway WebSocket, Bedrock, DynamoDB, S3, SNS, WAF, EC2, IAM, CloudWatch, CloudTrail, CloudFormation), Serverless Architecture, Azure, Docker, CI/CD (Jenkins), Linux
**Operations:** Production Monitoring, Root Cause Analysis, Incident Management

## PROFESSIONAL EXPERIENCE

**AI Engineer**
**Stellantis | USA | Aug 2026 - Present**

- Automated service-parts setup for Mexico vehicle launches in Python, pulling global parts data from Parts Smart, EV platform databases, and LCVD to identify launch-specific components and replace manual multi-country filtering.
- Built deterministic validation pipelines that cross-check parts across LCVD, CSPS, DealerCONNECT, Service Box, Web CN, and STV for VIN applicability, sales codes, engineering attributes, and powertrain compatibility.
- Automated the EBOM-to-SBOM parts release workflow by generating structured parts-list extracts and triggering backend batch jobs, feeding procurement and supply-chain tracking.
- Designed an exception-detection workflow from analysis of failed releases that flags sales-code mismatches, engineering-flag inconsistencies, and BOM errors before SBOM release, reducing manual troubleshooting.

**AI Engineer**
**Cogent Integrated Business Solutions, Inc. | USA | Jun 2025 – Aug 2026**

- Designed the architecture for and built a production AI lead-generation agent on cogentibs.com using AWS serverless (API Gateway WebSocket, Lambda, Amazon Bedrock, DynamoDB, S3, SNS), engineered for a 3-second response time, 99.5% uptime, and $5–50 per month in AWS cost.
- Built LLM intent classification that identifies 4 visitor types (corporate client, consultant, partner, general) within the first 3 exchanges and routes each into its own conversation flow, answering from an approved knowledge base and handing off to a human when it cannot answer.
- Implemented a weighted lead-scoring engine over 8 conversation signals (business email, company, service interest, consent, and more) that sorts leads into High (70+), Medium (40–69), and Low tiers and routes email alerts to sales leadership by tier and visitor type.
- Integrated HubSpot CRM to create or update contacts with 12 lead properties, including score, tier, area of interest, and an LLM-generated conversation summary, using email-based duplicate detection to update existing records instead of duplicating them.
- Secured the agent with AWS WAF, least-privilege IAM, and encryption in transit and at rest; enforced GDPR-aligned PII minimization, guardrails that block confidential partner names, and audit logs of every conversation (session ID, timestamp, visitor type, lead score).
- Built a source-grounded RAG assistant over internal policies and manuals, and evaluated every chunking, hybrid-search, and reranking change on Recall@5 and MRR, improving retrieval accuracy by 28%.

**Cloud Support Engineer**
**HCL Technologies | India | Jan 2022 – Aug 2023**

- Supported AWS and hybrid production infrastructure for 100+ pharmaceutical and insurance applications, resolving outages on EC2, Linux, and Windows using CloudWatch metrics and system logs.
- Resolved S3 access and permission failures by tracing IAM roles, bucket policies, and CloudTrail logs while keeping security controls intact.
- Automated server health checks and deployment support with Jenkins, CloudFormation, and Bash, reducing manual operational effort by 25%.
- Led L1/L2 incident triage and root cause analysis with cross-functional teams, sustaining 98% SLA compliance.

## EDUCATION

**MS in Computer Science (GPA 3.8/4.0) | Aug 2023 – Apr 2025**
Florida Atlantic University, Boca Raton, FL

**B.Tech in Electronics & Communication Engineering (8.5/10) | Aug 2017 – Jul 2021**
JNTU Kakinada, India

## PROJECTS

**ApplyEdge AI – Resume & Job Application Platform | Live**

- Built and deployed an LLM app that scores resume-to-JD match, rewrites weak bullets, and generates cover letters and interview questions, with live job listings from the Adzuna REST API.

**QueryVault – RAG Retrieval Agent | GitHub**

- Built a RAG pipeline with URL/PDF ingestion, FAISS top-k retrieval, source-chunk citations, and hallucination guardrails over 200-page documents.